

The Design and Build of a Simple Personal Finance System, Focused on Budgeting and Expenditure Analysis

Claudius de Moura Brasil
BSc Computing Project Report
Birkbeck College, University of London

May 2018

This report is the result of my own work except where explicitly stated in the text. The report may be freely copied and distributed provided the source is explicitly acknowledged.

1 Abstract

Personal finance systems exist in abundance nowadays, from open source to proprietary ones. They all tend to revolve around a basic common theme: providing accurate information about an individual's income and expenditure. Beyond this, they tend to vary in which features are implemented. The system designed and built for this project focuses on the use of the bookkeeping principle of double entry and the concept of pattern matching to find effective ways to categorise a user's expenditure, and provide them with relevant financial information to assist in decision making.

Contents

1	Abstract	i
2	Introduction	1
2.1	Requirements Capture Methods	2
2.2	Analysis and Design	2
2.2.1	Analysis and Design Patterns	2
3	Requirements	4
3.1	Business Case	4
3.2	Functional Requirements	4
3.2.1	Use Case Diagram	5
3.2.2	Use Case List	5
3.2.3	Designing the User Experience with Wireframes	6
3.3	Non-Functional Requirements	9
4	Analysis and Design	10
4.1	Business Logic	10
4.1.1	Create Manual Entry	12
4.1.2	Upload Statement	14
4.1.3	Visualise Categorised Summary	15
4.1.4	Calculate Budget	15
4.1.5	Final Class Diagram	16
4.2	The Persistence Layer	18
4.3	Database Diagram	18

5	Implementation	19
5.1	Presentation Layer	19
5.2	Business Logic	21
5.2.1	Scala Case Classes and the Prototype Design Pattern	21
5.2.2	The engine behind it all: the <code>Classifier</code>	22
5.3	The Presentation Layer	24
5.3.1	Algebraic Data Types and the Value Object Pattern	24
6	Testing	26
7	Conclusions	27
8	Reflections	28
8.1	Use Case Templates	28
8.2	Nested iterations in Analysis and Design stage	28
8.3	Trade-off between less code duplication and more sub-package independence	28
8.4	Dependency Injection	29
8.5	Implementation of Presentation Layer using only Functional Paradigm . .	29
8.6	Layering vs Manual Dependency Injection	29
8.7	Not implemented due to time constraints	30
8.7.1	Better exception handling and communication with the user	30
8.7.2	Validation	30
8.7.3	Implementation of the Strategy Design Pattern for Parser	30
	References	33
9	Appendix I: Types of Requirements and the use of UML	34

9.1 The use of Universal Modelling Language (UML) constructs	34
10 Appendix II: Previous versions of analysis patterns	35
11 Appendix III: Not-implemented version of Category with a field for balance	37
12 Appendix IV: The Estimate Tax Feature	38
12.1 Analysing the Estimate Tax Requirement	40

2 Introduction

A system could be summarised roughly as a solution to one or more problems. One of the first steps in order to build this kind of solution is to try to understand the problem – that is, try to map the requirements of the software. Vaasen et al. (2009 cited Boczeko, 2012, p. 8) suggests that an accounting information system’s main purpose is to provide information to internal and external stakeholders. Although this refers to accounting systems for businesses, it could be argued that this same definition can be employed to define personal finance systems – except that, in this case, the main stakeholder would be the individual using the system (that is, the user). In fact, one of the most widely known accounting systems available in the market, Quicken™, was conceived around the idea that there should be more efficient and less tedious ways to organise one’s personal financial information than doing it manually (Quicken Inc., 2017). This project has been developed based on similar ideas.

It seems fair to infer that nowadays most of a user’s financial transactions happen in ways that can be listed electronically (usually via their bank or credit card statements) – a study by Payments UK (2017), for example, indicates that there has been a rise in debit card payments over the past few years, and that the volumes of this type of transaction is likely to be higher than that of cash payments by the year 2021. Therefore, an assumption has been made that the users will require means of uploading a list of their financial transactions into the system.

The system created for this project intends to do just this. Its main feature, however, will be to allow the user to categorise expenditure based on patterns in the entries’ descriptions. Aside from this, there will also be a feature to allow the user to view summaries of the income and expenditure over a period of time, and another one to generate budget forecasts for future periods based on the financial information already entered.

This report documents the work of the project. Each chapter delineates a specific aspect of the development life cycle, which is in line with the development process described in the following paragraphs. Chapter 3 outlines the identified requirements which were used as motivation for the system to be developed; the contents of Chapter 4 shows further analysis and concomitant design of the system and the solutions it brings; Chapter 5 outlines select aspects of the implementation stage which serve to emphasise the techniques used to implement the designed logic, or highlights areas where it was felt it was necessary to implement something slightly different than what was designed.

The system developed for this project has been modelled after the principle of *double entry bookkeeping*, from the accounting domain, which states that “money is never created or destroyed – it merely moves from one account to another” (Fowler, 1997, Section 6.2). More specifically, double entry is the principle which ensures that every transaction always affects two accounts, one being credited (Out) and one debited (In). An account, for the scope of this project, refers either to a category created by the user, or to the user’s *cash book* – the contents of their bank account plus any manual entry which they make. In bookkeeping, each account can be classified as *asset*, *liability*, *income* or *expenditure*. Whether the account increases or decreases will depend on which of these categories it

falls under: *debits* will increase *assets* and *expense* accounts, and *credits* will increase *liability*, *capital* or *income* accounts (Wood et al., 2004, pp. 18-19).

Regarding the development method, an approach similar to that adopted by Bennett et al. (2010, p. 77) regarding software analysis and design has been employed, where no specific named methodology is espoused, but concepts of object-oriented analysis and design were applied, in an iterative and incremental fashion, using UML. More details about which concepts were used and the methodologies which originated them can be found in the following subsections.

The remaining definitions from Appendix I, including those of functional and non-functional requirements, will be employed when trying to classify the requirements and model the problem domain. The initial iterations will be focused more on the functional and usability requirements, paying some attention as well to specific non-functional requirements such as performance and security.

2.1 Requirements Capture Methods

The closest match identified to the techniques utilised for requirements capture for this project was ‘*Knowledge Acquisition*’, this relates to the process of capturing knowledge from an expert (Bennett et al., 2010, p. 150). In this particular case, though perhaps not qualifying as an expert, the author’s own qualifications and experience in accounting and bookkeeping were used to generate the ideas from which the requirements were extracted.

Differently from the original plan, and more in line with the *incremental model* of development exemplified by Dawson, 2009, pp. 120-124, the requirements were gathered in full before any actual development started.

2.2 Analysis and Design

As described by Bennett et al (2010, p. 348), “in projects that follow an iterative life cycle, design is not such a clear-cut stage, but rather an activity that will be carried out on the evolving model of the system”. Seeing that the development method being followed in this project is based on what should be an iterative approach, it was decided that the analysis and design of it would be done concurrently.

2.2.1 Analysis and Design Patterns

Fowler (1997, Section 1.3) defines a pattern as “an idea that has been useful in one context and will probably be useful in others”. This project will therefore attempt to utilise patterns where appropriate in order to prove this concept, and as an attempt to make use of the experience already acquired in the domain (or domains) in question. As emphasised by Bennett et al. (2010, p. 252), “a pattern is useful when it captures the

essence of a problem and a possible solution, without being too prescriptive”. So it may be the case that some of the patterns will be modified where necessary to optimally solve a problem.

The term *domain* is being employed, as defined by Evans (2004, p. 2), to define the “activity or interest of its user” – the “subject area to which the user applies the program”. So, analysis patterns will be used “when trying to understand the problem” domain (Fowler, 1997, Section 1.1). Essentially, an analysis pattern consists of a structure of classes and associations which occurs often in many modelling situations related to specific domains (Bennett et al., 2010, p. 254).

3 Requirements

3.1 Business Case

Any personal accounting system should be able to provide accurate and relevant summaries of an individual's financial status. In order to do this, the user needs to be able to supply the system with the necessary data, so that it can be analysed and properly converted into knowledge.

The scope of the personal finance system developed for this project must include a feature to allow a user to upload their bank and credit card statements into it. Each transaction should be classified based on categories which the user will create. The user should then be able to visualise a summary of their income and expenditure by category and period, which should allow them to have a concise and clear visibility of how much they have earned and spent over any period of time. The category must be handled with care – there should not be a case where a user deletes a category and then all the entries in a period are lost with it, but if a user wants to change the name of a category they should be free to do so.

Since there may be other sources of income which the user may want to categorise (such as breakdowns of cash transactions from pocket money), a feature to allow for manual entries should also be made available. The user can declare a lump sum, and then break it down among categories. The option should allow them to choose whether the transaction is a credit or a debit from a specific category, and then provide the corresponding debit or credit to a category of their choice. For example, if the user withdraws £50, and spends half of it on weekly shopping and the other half on a cinema ticket, they should be able to 'credit' (withdraw) £50 from the *Bank* category, and then 'debit' £25 to both *Weekly Shopping* and *Entertainment* categories.

Once the system has enough data, it should be able to calculate a simple budget and display it for the user. The budget can be a simple average of income and expenditure over a long enough period of time, projected over the future month/year – it would only be used as a guideline for the user, anyway.

3.2 Functional Requirements

Based on the description above, a few functional requirements were identified. They are represented in the diagram on section 3.2.1 and the list, wireframes and activity diagrams on section 3.2.2.

3.2.1 Use Case Diagram

Use case diagrams are UML constructs which were developed by Jacobson et al. (1992, cited Bennett et al., 2010, p. 154). The use case diagram on Figure 1 is used to illustrate the functional requirements identified for this project:




Figure 1: Use Case Diagram

3.2.2 Use Case List

Table 1 lists the descriptions for the use cases listed above:

Use Case	Description
Upload Statement	The user must be able to upload a list of their financial transactions, most likely their bank or credit card statements, in a valid format, and all entries should be categorised based on specific patterns <i>Includes:</i> Categorise Entries
Create Manual Entry	The user should be able to create a manual transaction for income or expenditure, include a date, amount and description, and classify it among existing categories or create new ones in the process <i>Extends:</i> Create Category
Visualise Categorised Summary	The user must be able to visualise a summary of their income and expenditure over a period of time
Calculate Budget	The user must be able to visualise a budget for future periods based on their income and expenditure data already entered
Categorise Items	Analyse the current entry and assign it to a category <i>Extends:</i> Create Category
Create Category	Creates a new category with the name suggested by the user

Table 1: Use Case Descriptions

The *Estimate Tax* feature was not included in these requirements, as the time constraints would not allowed for it to be implemented, but its specifications can be seen in [Appendix IV](#).

3.2.3 Designing the User Experience with Wireframes

The wireframe below (Figure 2) was created to better illustrate the *Manual Entry* requirement from the point of view of the user. It shows an example of an entry for a laptop and a licence for a proprietary operating system, which can then be broken down among different categories. The user has the option to use the percentage or the amount boxes in order to provide a breakdown, and they can also add new lines if more than one is required – the example shows two lines, but the default would be one. Under the category search box, if the user types a category name that does not exist they will be asked if they want to create a new one:

Manual Entry

Type

Income/Expenditure ▼

Date

19/05/2018 ▼

Total

1000.00

Currency

GBP ▼

Description

New Laptop with proprietary OS licence

Breakdown

Category	%	Amount
Laptops	90%	900.00
(start typing for search suggestions)	10%	100.00

New Line

Subtotal: 1000.00

Cancel

Submit

Figure 2: User interface wireframe for *Create Manual Entry* use case

And in order to better understand the relationship between *Upload Statement* and *Categorise Entries*, the activity diagram below (Figure 3) was developed:




Figure 3

At the first few iterations, in order to provide a minimum viable product, the process of categorisation will be a blocking one consisting of multiple calls being made to the process above – for each call, the process will block awaiting user input when a category is not found. However, there are plans for future iterations where this process could be optimised by concurrency, and if there is enough time a more appropriate interface will be built for such a case.

The initial interface for uploading a statement should be a simple one, such as the one below (4), and at least initially the interface for manual entry will be used whenever user input is required:

Upload Statement

CSV Path

(enter absolute file path or click 'Find')

Find

Upload

Please choose the csv version of the statement, then press 'Upload'

Figure 4

Below (Figure 5) is also a wireframe illustrating the GUI for *Visualise Categorised Summary*. The user should have an option to select the dates and, if they only want to see one category, the category itself:

Visualise Summary

From: To:

Start 19/05/2018

Category (optional):

(start typing to initiate search)

Cancel Submit

Figure 5

The dates field should allow them both to type a date or choose it from a drop down calendar. If no values are entered, the system will return a summary of all categories on the system.

3.3 Non-Functional Requirements

At this iteration, no significant non-functional requirements were identified, so they were not included in this report.

4 Analysis and Design

The system will be designed in a layered architecture, and will consist of three layers: Presentation, Business Logic and Persistence. The presentation layer will take care of the vast majority of the interactions between the user and the system, and will depend only on the Business Logic layer below it. The business logic will contain the model of the concepts which will attempt to implement a solution to the problems outlined in the requirements, and will depend only on the persistence layer below it. Lastly, the persistence layer will ensure that the data is not lost when the application is stopped (Bauer et al., 2016, p. 32-33). This approach is similar also to what has been described by Holmes (2016, p. 3) as full stack development for an application using a Model-View-Controller (MVC) architectural pattern.

Most of the aesthetic designs of the presentation layer happened in chapter 3.2, so this one will focus on the analysis and design of the business logic and persistence ones.

4.1 Business Logic

Seeing that this is a system dealing with finance, it would make sense to treat the categories as if they were treated as *accounts*. So, in order to make sure to imbue this system with knowledge acquired by more experienced programmers, analysis and design patterns will be employed.

It is also useful at this point to make a distinction between the types of classes used to model the domain between three possible kinds: the first are the classes which model the interaction between the system and its actors – these are called *boundary classes*; the second kind are those classes which model information and/or behaviour or some concept or phenomenon – these will be called *entity classes*; and lastly, there are those classes which model transactions, coordination, control and sequencing of other objects – which are known as *control classes* (Jacobson et al., 1999, cited Bennett et al., 2010, pp. 198-201).

The first analysis patterns which seem appropriate are a modified version of the *Account* pattern, used to create the **Category** entity class, and the *Quantity* pattern for the **Amount** entity class (Fowler, 1997, Sections 6.1 & 3.1):




Figure 6

As implied by the diagram above, the **Category** class will be associated with instances of the **Entry** class. This is done so that the only way to change the total of a category is by adding positive or negative entries to it – for example, to indicate a credit to a category, a negative entry can be added to it. One of the modifications to the original Account pattern consists of the fact that, whereas in the original pattern an instance of **Account** would keep track of the balance, there is no need to keep track of the current balance in each **Category** – the purpose of the system is to allow the user to view a summary of income/expenditure by period, so this will have to be calculated each time.

Another design choice which can be observed in Figure 15 is that the **Amount** class also possesses an attribute for currency. This has been designed so as to allow for the possibility of extending the system to keep track of transactions in multiple currencies, although this was not a specific requirements. Initially, there will only be a single default currency (GBP), so this will be hardcoded. But the class should be implemented in a way which also allows for different currencies to be loaded from another layer, such as the database, or an external API.

The next step is to provide a way for these entries to be added to categories. For this to happen, there needs to be a constraint to ensure that double entry happens every time a change needs to be made to a category. One of the ways to achieve this is to apply the *Transaction* pattern (Fowler, 1997, Section 6.2):




Figure 7

After having determined some of the analysis patterns which shall be employed, it makes sense to dive into a more specific analysis of the use cases described in Chapter 3. At this point the objective will be to start modelling classes and interactions based on concepts or things found in the problem domain. This will be done in the following subsections.

4.1.1 Create Manual Entry

The *Create Manual Entry* use case, which allows a user to input financial transactions individually using a specific interface, can be modelled as follows (Figure 9):




Figure 8: Communication Diagram for the *Create Manual Entry* use case

As the diagram above indicates, the **Transaction** class is responsible for the creation of new instances of **Entry** and **Amount**, which then get assigned to the **Category** instances chosen, or created, by the user. Once the user starts typing, the **CategorySuggester** is triggered to suggest categories. The actual implementation of how this suggestions happen may vary, but initially it should at least be based on what the user types in the search box. If the user chooses to create a new category instead, then they will be taken to the appropriate interface to allow them to do so. Once the user is satisfied and chooses to submit, the system will start to input the transactions in the appropriate category or categories.

Although **Transaction** was originally thought of as a *control class*, it was later decided that it would be beneficial to save the transaction information, so it was made into an *entity class* instead. It is also worth noting that, although implicit in the diagram and the wireframe on Figure 2, it is estimated that the user will start the interface, enter the transaction's details, and then start the process to find a category. Seeing that there will be a search suggestion at this point, it felt it was appropriate to have the diagram on Figure 9 have its first message sent at this stage.

It is also important to emphasise the fact that **CategorySuggester** is only an interface at this point, and that the suggerer in the diagram will be any object which

implements this interface. This is to allow more flexibility in the implementation of the classes responsible for suggesting categories to the user.

4.1.2 Upload Statement

The user should be able to upload their bank statements, provided that they are in a suitable format. The specifics of the format will be described in the implementation phase, together with more information on how to encapsulate as much as possible the complexities related to the formatting of this information. For the analysis and design phases, the emphasis will be on modelling the objects and their interactions. The diagram (Figure 9) below illustrates this process:

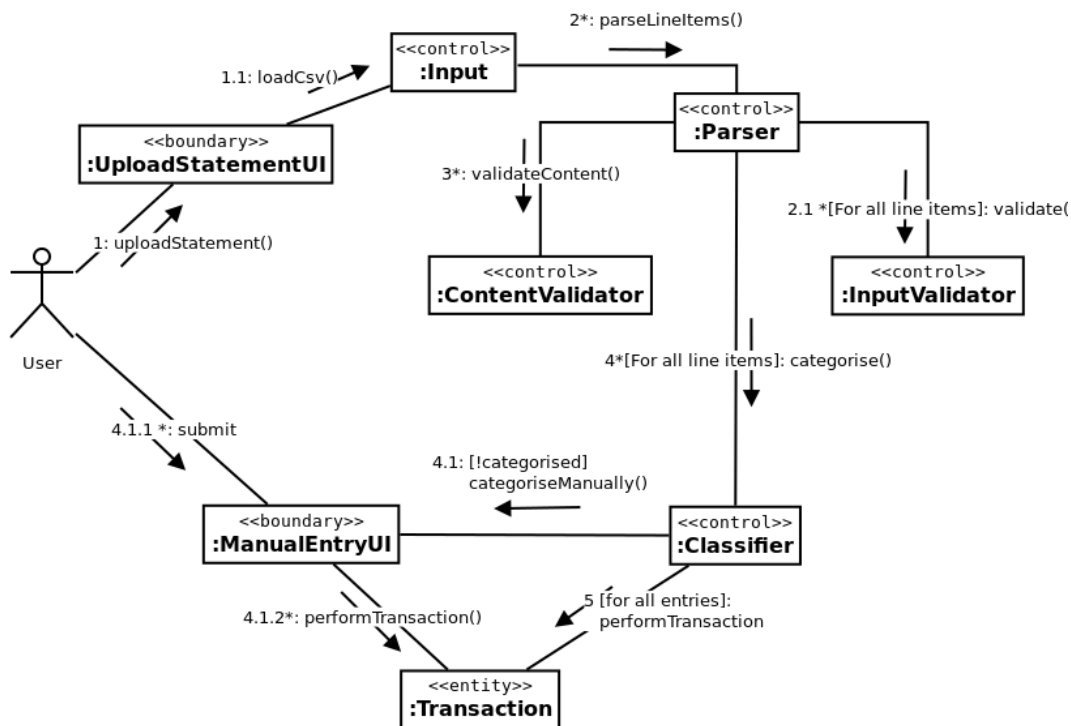


Figure 9: Communication Diagram for the *Upload Statement* use case

As can be seen in the diagram above, the process is spread among many classes. After the user uploads the statement, the loaded raw input will be sent to a **Parser** which will separate it according to the columns into the appropriate fields and line items. Then, what is now a collection of statement line items will be passed on to an **InputValidator** to make sure the user input is valid. Lastly, the resulting validated entries will be sent down to a **Classifier**, which will signal the relevant **Transaction** instance(s) to add the entries to their relevant **Category** instances – this last part has already been illustrated in Figure 9.

When the **Classifier** cannot match a line item against any of the existing categories, it will pass the line in question to the **ManualEntryUI**, which, apart from having all transaction details already populated, will rely on the process described at subsection

4.1.1 to properly categorise the item and then forward it to **Transaction** to create the categories.

4.1.3 Visualise Categorised Summary

As mentioned before, this feature will allow the user to view a summary of their income and expenditure by category. The user will enter the period which they want to examine, and the system will retrieve the categories which have entries with those dates. The system should then sort the categories by income or expenditure and by total, and then display it to the user. Optionally, the user can filter the output further by choosing a single category.

TODO: see Bennett et al., 2010, p. 262: sequence diagrams are for simple interactions. Move this to one of those

Below (Figure 10) is a communication diagram to illustrate this interaction:




Figure 10: communication diagram for *Visualise Categorised Summary* use case

4.1.4 Calculate Budget

This feature will allow the user to request the system to calculate a budget for them over a period of time. What the system will do then is retrieve all categories over the last 12 months, then calculate their means and take them as ratio of the income over the same last period. Then, it will take the first few which make up more than 80% of the total, and add all the others which make up the remainder and add them up as ‘other’ or something similar. It will then show these totals to the user, in a similar way as it shows the categorised summary. The communication diagram below (Figure 11) illustrates the steps taken by the application layer.

TODO: see Bennett et al., 2010, p. 262: sequence diagrams are for simple interactions. Consider moving this to one of those if the transaction can be simplified




Figure 11: communication diagram for *Calculate Budget*

4.1.5 Final Class Diagram

The class diagram on Figure 13 is the result of joining the more relevant entities listed so far:




Figure 12: Class diagram of the most relevant classes of the Business Logic Layer

Manual entries and uploaded statements will inevitably always go against a Bank category. For this reason, there does not seem to be any harm in hard-coding the **Bank** category. In this case, since most transactions are going to affect it, it was decided that no patterns should be assigned to it. Due to this, it was decided to allow multiplicity of `0..*` to the association between **Pattern** and **Category**. This relationship has also been made into a composition, since patterns should not exist without categories on the system, but a category can exist without a pattern.

4.2 The Persistence Layer

This is the layer which communicates directly with the database. It tries to encapsulate as much as possible the details of how it accesses the outside world by only exposing one singleton object, the `PersistenceMediator`. This is the object responsible for manipulating the database, with the help of other encapsulated package members. It uses a *Java Properties* file to choose which database engine to use (at the moment only *MySQL* and *H2* are available), but

4.3 Database Diagram

The UML diagram below shows the design of the database schema for the application:




Figure 13: UML Diagram of database schema

As the diagram indicates, the entities and relationships of the database are very similar to those of the business logic layer, except that it was felt that, since due to double-entry the `Entry` instances will occur multiple time, the description should be given a table of its own in order to contribute to optimising storage space usage.

One of the main reasons for choosing to model the application logic first, and then model the database, was that it would be more likely to ensure that the model (data structure) would reflect the needs of the application more accurately, as stated by Holmes (2018, p. 141). Also according to Holmes, the risk of having the model of the data adversely affect how the application will work and behave can be higher, if the application design starts with the data model.

5 Implementation

The initial Minimum Viable Product (MVP) was implemented with major differences from the design outlined in chapter 4. The most visible one is the fact that (what is mainly) the *Mediator* pattern (Nikolov, 2016, Ch. 9, Location 3594) has been used in order to implement the dialogue between the presentation and other layers using a feature which resembles the MVC architecture. Normally, this type of architecture is implemented in order to make sure that multiple views can refer to the same core functionality (or model) without actually implementing that functionality themselves, which allows the view to be responsible for presentation and the model for liaising with the business logic through a controller. This also allows the controller to update multiple existing views based on what input was received from one view (Bennett et al., 2010, p. 381). The difference in the implementation for this project is that this last feature, essentially, was not implemented, since only one view exists at one time.

In the implementation phase of this project, the components were segregated into packages (each of which represents one of the layers mentioned in Chapter 4) and sub-packages. Each sub-package has been implemented so that it does not have any knowledge of the super-package, but rather exposes and depends on interfaces which exist within itself. The super-packages, then, implement the interfaces of the sub-packages, which allows them to interact with the sub-package components. The *Mediator* pattern was implemented at the default package level by the *InteractionMediator* and *PersistenceMediator* objects, both of which also happen to represent Scala’s native implementation of the *Singleton* pattern (Nikolov, 2016, Ch. 6, Location. 2242). The following subsections will give more detail into the inner workings of the application and the techniques used to build it.

5.1 Presentation Layer

In the presentation layer, for this MVP, the *Scala Swing* package was used to design the *view*. The package itself consists of Scala wrappers for the *Java Swing* package, and one of the reasons why it was chosen was that, as with many other GUI packages, it already comes with an implementation of the *Observer Design Pattern* (Gamma et al., 1995, p. 293) in its capacity to react to events (Maier, 2009, p. 5; Nikolov, 2016, Ch. 9, Location 3731).

The entry point of the application is the `App` object, in the main package. This application works as a main method, and its only task is to call the `InteractionMediator.startup` method, which will send a message through the `PresentationMediator` instance returned by the `PresentationBuilder` to start the view.

In its current implementation, the view is started by the `MainWindow` class, which extends the `scala.swing.MainFrame` class. A `MainFrame` is a `Swing Frame` – “a window containing arbitrary data” (Odersky et al., 2016, Ch. 34, Section 34.1) – that can send a signal to terminate the application when closed. `MainWindow` also has implementations of other methods to allow the application to close gracefully, all of which were inspired by the `SimpleSwingApplication` of the `scala` package (Maier, 2009, p. 2 & 3). The author

chose not to make the first extend the latter was because the super-class implements a main method, which, considering that it would be another entry point to the application, could cause confusion when it were to start (not to mention the fact that this is not the responsibility of the `MainWindow` class).

As mentioned in the beginning of this chapter, each sub-package has been designed so that it does not have to depend on the packages which envelop it. The diagram below (Figure 14) illustrates this using the interaction between the `InteractionMediator`, `SwingAmbassador` and the `presentation.swing` package and sub-packages as an example:




Figure 14

In order to preserve encapsulation, specifically for the `swing` package, only the `MainWindow` controller and `SwingMediator` interface are exposed to the outside world, with the first having a constructor dependency on the latter. This means that, in order to interact with the components of the package, any external classes will need to implement the `SwingMediator` interface. This is also useful if the `presentation.swing` package is to be extracted and used as a library in another project, and also leaves room for the *Adaptor* pattern (Gamma et al., 1995, p. 139), if it needs to interact with other objects which do not necessarily implement the interface upon which it depends.

One feature also noticeable in Figure 14 is the fact the `personalfinance` package is unaware of which of the view's implementation it is interacting with. The modified version of *Builder* pattern (Løkke, 2009, pp. 37-38) implemented on `PresentationBuilder`, which uses reflection to determine what implementation to build, allows for the view itself to be decided with the use of an external `.properties` file. This allows for the actual implementation which will be delivered by the Factory to change without the code having to be recompiled, which introduces flexibility to the application. The `play` package is only a placeholder at this point, as due to time constraints it was not possible to actually

implement it, but it is there to exemplify how the project can be further extended.

Another element not seen in the diagram but which is present in the code is the fact that the dependency of the `MainWindow` class on the `SwingMediator` interface exists in the constructor, as seen in the code listing below:

```
class MainWindow(mediator: SwingMediator) extends Reactor {
```

This is an attempt at implementing the principle of *Dependency Injection* through *Constructor Injection* (Fowler, 2004). In the case of this project, the advantage of this simple piece of code is tremendous, especially because the fact that the `MainWindow` class does not have to instantiate a mediator – it simply relies on the fields published by the `SwingMediator` trait. Likewise, the fact that the `InteractionMediator` singleton object has a dependency on the `PresentationBuilder` only, allows for the actual implementation of the view to change without it having to change.

Talk how the above is a risky move, since Scala’s mixin composition has this habit of picking up the implementation of the last trait mixed into it, so in this case you get `PresentationMediator` inheriting from `SwingMediator` and then `PlayMediator` – although both are traits, this could cause problems in the future if the `Play` one implements something which `Swing` doesn’t, so basically these traits must always be the exact same

talk about the reason why the presentation mediator’s methods returning `Unit` being because this way they can stay as implementation-agnostic regarding the other layers as they possibly can, again thinking that the interface needs to be able to change without the code being recompiled

5.2 Business Logic

The business logic layer is the one which adhered more firmly to the model designed in section 4.1.5. The following are a few highlights from the implementation which are worth mentioning.

5.2.1 Scala Case Classes and the Prototype Design Pattern

One of the benefits of Scala are its case classes. Not just are they very useful for pattern matching, but also come with a few perks such as the `copy` method. This method allows for an object to be copied with some or all its members modified. It could be said that it is a language-native implementation of the *Prototype Design Pattern* (Nikolov, 2016, Ch. 6, Location 2461), and throughout the implementation and testing stage it proved to be a useful asset.

One of its many utilisations can be seen in the `Transaction` class, as one of the tools used to add entries to a `Category` without having to change the state of a specific instance – more similar to what is done in the *Functional Programming* paradigm:

Listing 2: extract of the `Transaction` class showing the `copy()` method in action

```
private def addEntries(tu: TransactionUnit): Category = {  
  val cat: Category = tu.category  
  val ents: Seq[Entry] = tu.entries  
  if (cat.entries.isEmpty) cat.copy(entries = ents)  
  else cat.copy(entries = cat.entries ++ ents)
```

TODO: write about why `Validator` has an auxiliary method for type – generics would affect the interface. The `Validator` was created as a trait in the first place to allow for lesser coupling between the classes which need validation

In order to be concise where possible, where it was felt that the domain’s requirements were captured accurately and was estimated that they are not likely to change, high level modules are depending directly on classes rather than interfaces. This can be seen as being in direct violation of the *Dependency Inversion* principle of SOLID (Martin, 1996), which some would consider to be a “risky” move (after all, it is not always possible to predict which requirements will and which will not change), but others would believe that this was acceptable for more static requirements. The author decided to follow the latter.

However, this principle was still employed where it seemed beneficial to follow it – for example, in areas where it was felt that the implementation of certain algorithms could be further optimised in future iterations. This is the case with `transaction.RegexDateStringParser`: this class implements the `transaction.dates.DateStringParser` trait, upon which the higher level modules which use dates will depend. This is an example of one of the places in the code where the *Dependency Inversion* principle can be seen, but it is not the only one.

5.2.2 The engine behind it all: the Classifier

Another example of *Dependency Inversion* (though not of *Dependency Injection*) is the `Classifier` trait, and the class implementing it: `StringClassifier`. The `InteractionMediator` declares a variable of the first type, and then instantiates it with the latter. Therefore, although it still depends on the abstraction, it also has a dependency on the implementation, which would mean that in order to change it a slight refactoring would be necessary.

For this implementation, `StringClassifier` was built using longest matching prefix to match entries to categories: patterns are matched to the start of the description of every entry, starting from longest pattern to the shortest. This is so that if there is, for example, a pattern “Laptop for parent” which would match to category “Gifts”, and another pattern “laptop”, for category “personal equipment”, then entries with the description “Laptop for Girlfriend” would not be picked up by the shorter match for laptop. Below is a code listing which shows how this was implemented:

Listing 3: Extract of the `StringClassifier` class, showing the algorithm currently used to match Entry descriptions’ prefixes against patterns

```

private def classifyByDescription(categories: Seq[Category],
    entries: Seq[Entry]): (Seq[TransactionUnit], Seq[Entry]) = {

    // extract all the patterns from categories and sort
    // them by length, from longest to shortest
    val patternIndex: Seq[String] = categories
        .flatMap(c=>{c.patterns.list.map( _.value )})
        .sortBy(_.length)
        .reverse

    // match the full pattern against the prefix of every
    // entry description
    entries.map({
        entry => (entry, patternIndex.find(entry.description.startsWith)) match {

            // if it finds them, return a transaction unit for
            // that category
            case (e: Entry, Some(pat)) =>
                TransactionUnit(categories.find(cat => {
                    cat.patterns.list.foldLeft(false)(
                        (test, catPat) =>
                            catPat.value.toLowerCase() == pat.toLowerCase() || test
                    ))
                ).get, Seq(e))

            // if not, return the entry as uncategorised
            case (e: Entry, None) => e
        }
    }).partition(_.isInstanceOf[TransactionUnit])
        .asInstanceOf[(Seq[TransactionUnit], Seq[Entry])]
}

```

As mentioned on the code listing above, the *longest matching prefix* is achieved by sorting all the patterns by length, from longest to shortest, then trying to match them against the prefix of the entry description in the same order. This way, even if an entry would be a match against more than one category, the pattern with the longest matching prefix will still be the first match.

One other interesting feature which can be seen in the above listing is how much is being done using the native `flatMap` and `foldLeft` methods of Scala collections. The `map` function and its derivatives are normally implemented by types which can be classified as *Functors* in scala. Functors are challenging to explain concisely, but for this report it should suffice to say that they are *classes which implement the map method and conform to a set of laws called functor laws*. The `foldLeft` method of the collections are also an indication that they can be classified as *monoids*, which are another construct brought over from mathematics. The main point of mentioning these is because they are one of the things which makes the functional paradigm so powerful: these constructs create a common interface which allows for the different types with which they are associated to be interacted with in a common way – that is, a lot can be done with them with very few lines of code (Nikolov, 2016, Ch. 1, Location 4243 & 703). Not just this, but the immutability

of the instances within them also makes it easy to perform these operations in parallel and, although this parallelism has not been explicitly implemented in this iteration, the author felt it was worth mentioning them in this section.

5.3 The Presentation Layer

The presentation layer is another one where decoupling and reflection can be observed. The package is accessible via the `PersistenceBridge` class, which uses reflection to instantiate one of the classes which implements the `ConnectionType` trait. This is an implementation of the *Strategy* pattern (Løkke, 2009, p. 80), which allows for the database vendor to be changed without the code having to be recompiled. At the moment, the only dialect implemented is *MySQL*, but a placeholder can already be seen for the *H2* database.

It is also worth noticing that the operations currently being executed by the implementation have a lot of room for improvement: many of the queries and updates which could be grouped together and sent to the database in batch are still being done individually. This does not have a huge impact on performance with a local database, but might cause issues with remote ones. Further iterations of the application should focus on optimising this.

Another pattern which can be observed in this layer is the *Value Object* pattern, which will be discussed in more detail below.

5.3.1 Algebraic Data Types and the Value Object Pattern

Throughout the code, examples of Algebraic Data Types can be seen. These appear in the form of sealed traits and case objects, and are normally used when instances need to be passed around as values, but also contain information which will be relevant to the code. Examples of these can be found in the `ConnectionType` hierarchy, where the number of possible instances for each case class would classify the trait and its subtypes as *Product Types* (Wampler et al., 2015, p. 411). The code listing below illustrates this:

Listing 4: extract of the `ConnectionType` hierarchy showing the case classes used as value objects

```
private[persistence] sealed trait ConnectionType {  
  
    ...  
  
private[persistence] final case class MySql(_dbName: String)  
    extends ConnectionType {  
  
    ...  
  
private[persistence] final case class H2(_dbName: String)  
    extends ConnectionType {
```

...

Algebraic data types could be said to be the natural implementation of the Value Object design pattern. This pattern is used widely for comparison of objects not by their identities, but rather by their values. They consist of small, immutable objects, and the instances of the case classes can be classified as just those (Nikolov, [2016](#), Ch. 8, Location 3068).

6 Testing

TODO: write about Scala's own implementation of Mockito as one of the language features

TODO: write about why I chose to do integration tests rather than unit

TODO: write about why I chose BDD as my form of TDD, if there's space write about the BehaviourTester class. Include Scala's infix notation as one of the reasons

TODO: write about how there is very limited testing in the presentation layer due to the fact that I do not have much experience with testing front end

TODO: talk about how testing for the persistence layer was done with MySql because the present version uses a local db, but that in the future for performance reasons they should be done using a self-generated portable db version instead (such as H2)

TODO: had to make a decision when testing the persistence regarding the reloading of the whole DB from the dump after performing Create/Update/Delete only. Reloading it every time using 'before' was just taking too long, so to improve performance I redid it in an unelegant way. The alternative of having one test suite for reading and another one for writing was not working also, as many problems were happening, possibly due to so many classes overwriting the database concurrently

7 Conclusions

Although not with as many iterations as originally planned, the main goal of the project has been achieved – that is, to deliver a functional, simple expenditure analyser. The techniques used to develop it should allow for the code to be easily extended and refactored as appropriate, without requiring too much of a rewrite.

The Scala language’s multi-paradigm features allowed for most of the variables in the code to be made immutable where appropriate, while still ensuring that the mutable aspects inherited from other languages worked as expected. The functional aspects of the language made it easy for the developer to work well with the immutability, and the flexibility of the data structures also allowed for the focus to be placed on writing good abstractions, while trusting that the language’s own implementation of its data structures, especially its collections (Odersky et al., 2016, Ch. 24), would maximise efficiency (as much as a possible).

Regarding the author’s perception of *software development best practices*, the advantages of following an iterative cycle were definitely visible, especially those of trying to model the system before beginning the implementation using an appropriate language (in this case, UML). The models in chapter 4 gave very good visibility of possible ‘pitfalls’, such as too much responsibility being assigned to an entity, or which level of abstraction should be used for each domain aspect being modelled. This allowed for them to be redesigned more easily than would be possible if no planning had taken place.

The benefits of analysis and design patterns were also noticed, and due to these techniques and the SOLID principles (which, although not applied perfectly, seem to have been used up to a reasonable extent), the code can now be extended more easily, and more features can be added with minimum impact being made to existing functionality. This shows, in the author’s opinion, the value of these techniques for software development, and why they are considered best practices.

8 Reflections

8.1 Use Case Templates

Originally, no template was used to document the use cases. The intention was to provide better ones at a later iteration, perhaps by researching the ones mentioned by Bennett et al. (2010, p. 157), but unfortunately there was not enough time, so the little there was had to be dedicated to the software itself.

8.2 Nested iterations in Analysis and Design stage

The Analysis and Design stage of the first iteration was delayed due to multiple ‘trial and errors’ within it. Appendices 10 and 11 show a examples of different models which had been incorporated into the final design, but which were later changed even before the implementation started – that is, even before any coding was done. This caused a reflection on whether the development method was truly iterative, and whether or not it was more similar to the Waterfall model.

Still, in the author’s opinion, spending more time within the analysis and design stage were very helpful in implementing *SOLID* classes, and as a result decreased the negative impacts that any refactoring in the code base would have caused, were it not to have been done as such. In previous iterations of similar projects by the author, but where the modelling normally done in the analysis/design stages were neglected, classes ended up with too many responsibilities and hard coded dependencies, which made any refactoring very challenging.

8.3 Trade-off between less code duplication and more sub-package independence

In order to keep sub-packages independent, as described in sub-section 5.1, a decision had to be made about increasing code duplication. As can be seen in the current implementation, both `PlayMediator` and `SwingMediator` are supposed to declare the same method signatures, so that there is no conflict when switching between implementations at run time. In order for this to happen, each interface had to be manually typed into its package, and the same will have to happen for any further extensions to the code. An alternative to this would be to allow the sub-packages containing the implementations to see the super-package, and declare the methods once in the `PresentationMediator` interface, but this would have meant sacrificing the concept of sub-package independence already mentioned.

8.4 Dependency Injection

Very often during the implementation it was felt that better dependency injection could be achieved. In classes such as those found in the `TransactionsValidators.scala` file there was constant hard-coded dependency to the `validation` package. In this instance, a framework such as *Spring* or *Guice* would have been useful, but unfortunately time constraints made it unlikely that the author would be able to successfully implement these and still deliver an MVP for this project. As a result, there is more tight coupling than there could be, but wherever possible an effort has been made to pass the dependencies as constructor or method parameters, so as to facilitate testing, among other things.

8.5 Implementation of Presentation Layer using only Functional Paradigm

THIS MAY HAVE BEEN CHANGED. LAST COMMIT WHERE TRUE WAS b0bae376264f After the first iteration of the presentation layer, it was noticed that perhaps it could be fully implemented in a form more close to that of the functional paradigm. That is, make it a point to not use `var`'s, and only `val`'s for members (also not change state using Scala Swing's classes natural mutable fields, such as location, visible, etc – the effects from these could be replicated by copying the values of the instances into new ones). This could have been achieved if the full `presentation.swing` package had been implemented from the start with this in mind: have an interaction mediator within the package, and then make more use of double dispatch, familiar to the *Visitor Pattern* (Nikolov, 2016, Ch. 8, Location 3943). Then, for the actual flow of the application, a strategy similar to that used by Felleisen et al. (2013, Ch. 5), where each action would trigger a function which changes the state of the application, and then the GUI displaying the new state would be passed recursively to the main function, ensures that the immutability of the functional paradigm is maintained throughout the presentation layer.

Unfortunately, time constraints were once again too tight for this to be fully implemented.

8.6 Layering vs Manual Dependency Injection

One of the original (implicit) goals of this project was to have a hierarchy system, where each sub-package would not depend super package, but super packages could depend on sub ones. That is, a highly specialised `presentation.swing.frames` package would not depend on elements of the `presentation` package, but the `swing.frames` package would declare interfaces which would then be implemented by `presentation`. What had not been taken into consideration, however, is how the fact that the super package having to implement the interface would make it difficult to truly implement dependency injection: the idea was to have the `InteractionMediator` implement the lower package's interfaces, so that it could be passed to the classes of the specialised package. But this would be a

problem when the interface is too specialised. Therefore, a compromise had to be made and the interfaces had to be made more generic.

8.7 Not implemented due to time constraints

The following topics were not implemented due to time constraints.

8.7.1 Better exception handling and communication with the user

The `PersistenceMediator` class should be handling exceptions which might be thrown by `PersistenceBridge`'s every time the first calls the latter's methods. The intention of allowing the Mediator to catch exceptions was because this could then be passed to the user as informative messages, or be handled internally depending on the nature of the exception (e.g., connection exceptions would have to be handled internally, and exceptions related to user input should be passed to the user). Therefore, some exceptions should be handled by the `PersistenceBridge`, and others by the mediator.

A feature to communicate with the end user was being implemented as well by means of the `presentation.swing.Messenger` object. Unfortunately, there were bugs in the code which did not allow it to work properly, and a solution could not be found in time for the deadline, which meant this feature was not implemented in the MVP.

8.7.2 Validation

A lot of thought has been put into where validation should happen. For example, the constraints of `Transaction`, `Category` and `Entry` which were used to enforce *double-entry* could have been implemented at database or application levels, or both. In the current implementation, however, this constraint is only enforced at business logic layer, and in the `PersistenceBridge.createEntrySet` method of the persistence layer.

8.7.3 Implementation of the Strategy Design Pattern for Parser

One of the original intentions of the author was to utilise an implementation of the Strategy Design pattern when loading the user's bank statements into the system. The current implementation uploads data from CSV files, but making use of the Strategy pattern, which allows for different implementations of an algorithm to exist, and for the right one to be chosen while the application is *running* (Nikolov, 2016, Ch. 8, Location 3152), would allow for other formats to be used too. JSON and XML formats come to mind, especially if a version of the application could be made which would allow for it to communicate with a banking system's API – API's (especially RESTful) tend to favour these two formats.

TODO: talk about the small amount of input validation due to the first version being for a desktop app only, but that if a hosted version were to be made available more validation would be added

talk about the decision to sometimes use pattern matching with objects and at other times use enum classes (as with the objects which extend `KitName` in the presentation layer vs the `EntryType` in the root `personalfinance` package), and that although here they have been used interchangeably, a level of consistence would have been improved readability in a different situation

Regarding the fact that the same transaction can be repeated multiple times with manual entry, this was allowed to happen because it is not clear whether it is impossible for multiple transactions in a statement to have the exact same creation date, description and value. The same is true for manual entries: if the user wants, they can enter as many entries with the same amount as they want

The lack of experience of the author really showed when it came to getting individual categories from the user. The idea of making subpackages unaware of super packages led to “awkward” uses of tuples and maps as parameters on the user interface, and a somewhat unelegant way of handling the remaining uncategorised entries after a new category has been created: since every method has a `Unit` value, the whole block of entries had to be passed around as messages so that the state would not be lost. Hence why there are so many `Seq String String String String` flying around.

Perhaps due to what is a combination of inexperience and lack of more detailed preparation regarding method signatures, there is a bit of (for lack of a better term) ‘awkwardness’ around method signatures, return types and constructors. For instance, `Transaction.execute` takes `Seq TransactionUnit`, but at times it is clear that only two transaction units are going to be entered, but it is still necessary to make them into a sequel. Or when `Classifier` makes a transaction unit, it still needs to wrap the category and entry into a `Seq` before it can create the method. Although it all works, there could be better ways to engineer it so that it is clear from the contract and the fields in the class what is happening inside of it

The way the matching of an entry description against a pattern is being done should be by: first sort all the entries by longest description; then categorise, create patterns etc; then when actually classifying similar entries in the next month, start by the longest transactions, and match them first against the longest patterns, then the shortest. This way more generic patterns such as ‘laptop’ would not match first with an entry such as ‘laptop for girlfriend’

Although originally there did not seem to be any advantages on implementing the ‘Amount’ field in the ‘Category’ class, as suggested by Fowler (1997) when referring to the ‘Account’ analysis pattern, it became apparent that it would have been beneficial to have it for the budgeting feature. If each

category could have their own amount, then instead of passing around Tuples of String and Double, instances of 'Category' could have been passed instead, each of which containing their own subtotal (as per [Appendix III](#)) during the period in question.

The current implementation of the budget feature would not work well if the user has overspent: it will increase the projected income to match the user's expenditure, which would not be a desirable trait in a commercial expenditure analyser. There are better ways to implement this feature, such as decreasing some of the expenditure when budgeting, but the time constraints once again prevented this from being built.

References

- Bauer, Christian, King, Gavin, and Gregory, Gary (2016). *Java Persistence with Hibernate*. Second Edition. Manning Publications Co.
- Bennett, Simon, McRobb, Steve, and Farmer, Ray (2010). *Object-oriented systems analysis and design using UML*. Fourth Edition. McGraw Hill Higher Education. ISBN: 978-0-07-712536-3.
- Boczko, Tony (2012). *Introduction to Accounting Information Systems*. Pearson Custom Publishing.
- Dawson, Christian W (2009). *Projects in computing and information systems: a student's guide*. Second Edition. Pearson Education.
- Evans, Eric (2004). *Domain-driven design: tackling complexity in the heart of software*. Addison-Wesley Professional.
- Felleisen, Matthias, Van Horn, David, Barski, Conrad, et al. (2013). *Realm of Racket: Learn to Program, One Game at a Time!* No Starch Press.
- Fowler, Martin (1997). *Analysis patterns: reusable object models (OBT)*. Kindle Edition. Addison-Wesley Professional.
- (2004). “Inversion of control containers and the dependency injection pattern”. In: Gamma, Erich, Helm, Richard, Johnson, Ralph, and Vlissides, John (1995). *Design patterns: elements of reusable object-oriented software*. Addison-Wesley.
- Holmes, Simon (2016). *Getting MEAN with Mongo, Express, Angular and Node*. Manning Publications Co.
- (2018). *Getting MEAN with Mongo, Express, Angular and Node*. Second Edition, Early Access, V7. Manning Publications Co.
- Løkke, Fredrik Skeel (2009). “Scala and Design Patterns-exploring Language Expressivity”. In: *Aarhus Universitet, Datalogisk Institut* 44.
- Maier, Ingo (2009). “The scala.swing package”. In: *Scala Improvement Process (SID)* 8.
- Martin, Robert C (1996). “The dependency inversion principle”. In: *C++ Report* 8.6, pp. 61–66.
- Nikolov, Ivan (2016). *Scala Design Patterns*. Kindle Edition. Packt Publishing Ltd.
- Odersky, Martin, Spoon, Lex, and Venners, Bill (2016). *Programming in Scala*. Third Edition, Kindle eBook. Artima Press.
- OMG (2015). *Unified Modeling Language Specification*. Version 2.5. URL: <http://www.omg.org/spec/UML/2.5/PDF> (visited on: 16 Dec. 2017).
- Payments UK Management Ltd. (2017). *UK Payments Markets Summary*. URL: <https://www.paymentsuk.org.uk/files/puk-uk-payment-markets-2017-summary.pdf> (visited on: 18 Dec. 2017).
- Quicken Inc. (2017). *About Us*. URL: <https://www.quicken.com/about-us> (visited on: 17 Dec. 2017).
- Wampler, Dean and Payne, Alex (2015). *Programming Scala*. Second Edition. “O’Reilly Media, Inc.”
- Wood, Frank and Robinson, Sheila I (2004). *Book-keeping and Accounts*. 6th Edition. Pearson Education.

9 Appendix I: Types of Requirements and the use of UML

Bennett et al. (2010, pp. 140-142) categorises requirements as being of three types:

Functional Requirements The system’s functionality – what it is expected to do.

Non-functional Requirements How well the system delivers its functionality. These requirements are related to the performance, scalability, availability, recovery time, security, and others.

Usability Requirements These relate to how effectively, efficiently and satisfactorily users can achieve their goals in the existing system. User interfaces can play a big part in meeting these requirements.

9.1 The use of Universal Modelling Language (UML) constructs

UML is a modelling language created with the intention of providing system architects, software engineers and developers with a common set of modelling tools, with a defined syntax, which would help them better analyse and design software-based systems, and to model business and similar processes (OMG, 2015, p. 43). It defines several constructs which have been employed throughout this report in order to model the specifications of the system, such as Use Case, Activity, Class, Sequence and Communication Diagrams.

In the class diagrams used, some thought was given to how to model the relationships between entities. In the end, it was decided that transitory relationships should be modelled as dependencies, and more permanent ones would be modelled as associations. An example can be seen in Figure 16, where the **Category** interface has an association to the **Entry** class, since entries are not going to be removed, only added, to categories, making it a more permanent relationship.

10 Appendix II: Previous versions of analysis patterns

Below is a previous version of analysis patterns. It was decided that keeping track of *balance* was not something worth pursuing, since the main point of the application is for users to choose a period of time and then look at a summary of expenditure in there.

The first analysis patterns which seem appropriate are the *Account* pattern, used to create the *Category* entity class, and the *Quantity* pattern for the *Amount* entity class (Fowler, 1997, Sections 6.1 & 3.1):





Figure 15

As implied by the diagram above, the *Category* class will be associated with instances of the *Entry* class, and will keep track of the balance made up of the sum of *Amount*'s of each entry. This is done so that the only way to change the total of a category is by adding positive or negative entries to it – for example, to indicate a credit to a category, a negative entry can be added to it.

Another design choice which can be observed in Figure 15 is that the *Amount* class also possesses an attribute for currency. This has been designed so as to allow for the possibility of extending the design to keep track of transactions in multiple currencies, although it was not a specific requirement. Initially, there will only be a single default currency which shall be set at runtime.

The next step is to provide a way for these entries to be added to categories. For this to happen, there needs to be a constraint to ensure that double entry happens every time a change needs to be made to a category. One of the ways to achieve this is to apply the *Transaction* pattern (Fowler, 1997, Section 6.2):



After having determined the analysis patterns which shall be employed, it makes sense to dive into a deeper analysis of the use cases described in Chapter 3. At this point the objective will be to start modelling classes based on concepts or things found in the problem domain. This will be done in the following subsections.

11 Appendix III: Not-implemented version of Category with a field for balance

The following text was part of a previous model of how to use the *Summary Account* pattern to try to classify accounts between different types. It has been since then decided that the best way would actually to have the categories themselves keep track of their own types.

Furthermore, especially due to the requirements for tax estimates, the *Summary Accounts* pattern (Fowler, 1997, Section 6.3) will be adapted to help classify categories between income and expenditure. At this point, it seemed sensible to make **Category** into an interface, and specialise it on its subtypes. A **SummaryCategory** implements the **Category** interface, and would implement its **getEntries** method so that its entries are those of its components. It's components are other instances of **Category**, so they can be both both **DetailCategory** and its own type – the implementation will have to take this into account somehow, such as by using recursion. The class diagram below (Figure 16) can illustrate it better:




Figure 16

12 Appendix IV: The Estimate Tax Feature

Use Case	Description
Estimate Tax	Based on the information entered and the current tax year, calculate how much tax is due

Table 2: Calculate Tax Requirement

A feature to estimate tax would also be useful. Initially, this can be based on the system used to calculate personal tax in the UK, where all of an individual's income is added up and, depending on which threshold it reaches, tax is deducted at that level. For example, for an individual earning £75,000 per annum from their job where he or she is a company employee, initially their tax free allowance would be deducted from the gross figure, then the basic rate is deducted from it up to its limit, then higher rate would be deducted from the rest. However, if they have any income from sources which have special taxation rules, such as interest on savings, these would be deducted right after the personal allowance – that is, income from special taxation sources get added to the gross income figure, but tax from them is deducted at a different rate and before the base tax. (CITATION NEEDED).

There should be an option to allow a user to determine whether a transaction is taxable. For income, the user should also be able to determine if it is already net of tax (taxed at source) or if it's the gross amount, since this would influence the tax calculation. Also, on expenditure, certain costs such as memberships to professional organisations are deductible, so the user should be able to indicate this when they are registering the category.

Furthermore, for the tax feature, it is important to emphasise that not all of an individual's income will get taxed under income tax. For example, profits on sales of shares gets taxed as capital gains tax. Since this type of tax is outside of the scope of this application but would still appear on a user's list of transactions, the system needs a way to highlight these so as for them not to be included in the tax calculation.

Below is a wireframe to illustrate how the initial window for the *Estimate Tax* requirement¹:

¹HM Revenue & Customs (2018). Income Tax rates and allowances: current and past. Url: <https://www.gov.uk/government/publications/rates-and-allowances-income-tax/income-tax-rates-and-allowances-current-and-past> (visited on: 03 Mar. 2018).

Estimate Tax

When does the tax year start? * ?
 06/Apr/2017 ▼

Tax Free Allowance
 £ 11,000.00

Special Taxation Rates - After Allowance ? **New Line**

Rate Name	% Deducted	Incremental Income ?
Savings	5%	£ 5,000.00

Normal Taxation Rates - After Special Rates ? **New Line**

Rate Name	% Deducted	Incremental Income ?
Basic Rate	20%	£ 31,500.00
Higher Rate	40%	£ 33,500.00

Back **Estimate**

Figure 17: Estimate Tax interface. Date and allowances shown as specified by HM Revenue & Customs

Normally, in the UK there is a difference between what is “tax exempt” and what is “taxed at 0%”. For the purposes of the estimation provided by the requirements in this project, it was initially decided that this distinction would be ignored, and that the user could simply have a choice to use one of the dynamic boxes on *Normal Taxation Rates* to indicate a tax free allowance. However, due to the order in which the taxes are to be deducted, it was decided that a non-taxable allowance field should still be made available, but allow its value to be set to zero. This should allow for more flexibility.

The interface would start with one line item, and then more items could be added if the end user needs them. Each line represents one of the tax rates and the threshold of income which needs to be reached before it starts deducting tax at that percentage.

A choice has been made to include help texts next to the fields which are likely to need them, and these are indicated in Figure 17 by the bubbles with the question marks next to the relevant fields.

12.1 Analysing the Estimate Tax Requirement

This feature will allow the user to calculate an estimate of the amount of tax due for a financial year, based on their income and expenditure. For this to happen, the user will have to provide the date when the financial year begins, and the tax tiers in their country. The tax estimation will be based as best as possible on how personal tax is calculated in the UK ².

It was felt that Fowler’s (1997, Session. 6.4-6.5) *Memo Account* and *Posting Rules* patterns would be useful for this case. Designing this requirement as such would ensure that, whenever the user created a category which was deemed taxable, a portion of the income would be posted to the tax account. It is also necessary to remember whether income has already been deducted or if it relates to the gross amount, so this will have to be taken into consideration.

The *Memo Account* pattern here will be implemented as tax-related instances of *Category*. Due to the double-entry principle of bookkeeping, each tax instance will need to have a *contra* category – that is, a second category which will act as the counterpart to the transaction – so that the principle is still observed. In order for the user not to have to create the transactions manually, however, an *event listener* needs to be added in order for entries to be made in the appropriate tax categories when necessary. This is where the *Posting Rules* pattern will be appropriate. Below (Figure 18) is an analysis class diagram exemplifying the implementation:




Figure 18

So the principle being illustrated here is that, whenever an entry is posted to a category which the user has marked as subscriber to tax, then the **PostingRule**, which subscribes to that event, would determine whether a transaction needs to happen between the tax category and its contra. Then, when estimates need to be generated, they can simply come directly from the tax categories.

²HM Revenue & Customs, HM Treasury, Department for Work and Pensions, and Others . Personal tax: Income Tax. Url: <https://www.gov.uk/topic/personal-tax/income-tax> (visited on: 03 Mar. 2018).